

Jared Joselowitz

📍 London, United Kingdom

✉ jaredjoss123@gmail.com

☎ 07440 408647

🌐 jaredjoselowitz

🔗 jaredjoss

SUMMARY

Lead AI Research Engineer owning the clinical safety-evaluation programme that gates every production release of LLM voice agents running across 200,000+ live patient calls. Leads a 4-person science team across safety-focused evaluation, multi-agent simulation, LLM-as-judge methods, and ASR robustness. First author at COLM 2025 on LLM interpretability through inverse reinforcement learning; co-author at ACL 2025 and IWSDS 2026.

EXPERIENCE

Ufonia

Feb 2025 – present

Lead AI Research Engineer (promoted from Senior, Apr 2026)

London, United Kingdom

- Own the clinical safety-evaluation programme that gates every production release of the LLM voice agents: define the validation standards and the simulated-patient and LLM-as-judge methods it runs on, provide the evidence base for clinical safety assurance, and set the team's research direction.
- Co-developed MATRIX, a multi-agent simulation framework for safety evaluation of clinical dialogue agents; its behaviour judge matched expert-clinician hazard detection (F1 0.96) in a blinded 240-dialogue assessment, and the framework benchmarked five LLM agents across 2,100 simulated dialogues, 14 hazard scenarios, and 10 clinical domains.
- Built the ASR-safety evaluation behind WER is Unaware, a clinician-annotated benchmark and a GEPA-optimised LLM-as-judge that reached 90% accuracy and Cohen's kappa 0.816 at matching clinician safety assessments and showed word error rate poorly reflects clinical risk.
- Built the clinical-dialogue evaluation pipelines behind these systems (reference-based, simulated-patient, and LLM-as-judge methods validated against clinician assessment), spanning hazards from ASR errors to patient interruptions and multilingual interactions.
- Benchmark open-source, multilingual, and speech-to-text models against the clinical safety suite to determine which are cleared for production.

The Awareness Company

Sept 2022 – July 2023

Junior Data Scientist

Johannesburg, South Africa

- Built computer-vision models (Azure Cognitive Services) and an energy-demand model, and designed, trained, and deployed a soil-moisture predictor and a poaching-risk model for Kafue National Park, Zambia.
- Used Azure OpenAI Service to fine-tune an LLM chatbot on company data, deployed to end users via an API.

Tesseract

Feb 2022 – Aug 2022

Junior Full-Stack Software Engineer

Johannesburg, South Africa

- Worked in a five-person team on a full-stack AI annotation platform for business intelligence in the FinTech industry, across frontend (React) and backend (Django).
- Built end-to-end testing with Cypress and a WhatsApp annotation bot deployed on AWS Lambda.

PUBLICATIONS

Insights from the Inverse: Reconstructing LLM Training Goals Through Inverse Reinforcement Learning

[arxiv.org](https://arxiv.org/abs/2505.10000)

Jared Joselowitz, Ritam Majumdar, Arjun Jagota, Matthieu Bou, Nyal Patel, Satyapriya Krishna, Sonali Parbhoo
Conference on Language Modeling (COLM 2025). [Code](#)

WER is Unaware: Assessing How ASR Errors Distort Clinical Understanding in Patient Facing Dialogue

[aclanthology.org](https://aclanthology.org/2025.acl-main.100/)

Zachary Ellis, Jared Joselowitz, Yash Deo, Yajie He, Anna Kalygina, Aisling Higham, Mana Rahimzadeh, Yan Jia, Ibrahim Habli, Ernest Lim

International Workshop on Spoken Dialogue Systems (IWSDS 2026), Oral Presentation. [Code](#)

ASTRID: An Automated and Scalable TRIaD for the Evaluation of RAG-based Clinical Question Answering Systems

[aclanthology.org](https://aclanthology.org/2025.acl-main.100/)

Yajie Vera He, Mohita Chowdhury, Jared Joselowitz, Aisling Higham, Ernest Lim
Findings of the Association for Computational Linguistics (ACL 2025)

MATRIX: Multi-Agent simulaTion fRamework for safe Interactions and conteXtual clinical conversational evaluation

[arxiv.org](https://arxiv.org/abs/2505.10000)

Ernest Lim, Yajie Vera He, Jared Joselowitz, Kate Preston, Mohita Chowdhury, Louis Williams, Aisling Higham, Katrina Mason, Mariane Melo, Tom Lawton, Yan Jia, Ibrahim Habli
arXiv preprint, under review (2025)

TALKS & PRESENTATIONS

Shipping AI to a Million Patients Without an A/B Test | *AI Engineer World's Fair 2026 (stage talk)*

- AI in Healthcare Track, San Francisco, June 2026.

When Patients Cut In: Extending Clinical Conversational AI Safety to Interruptions | *DAIH*

Workshop, COLM 2026 (oral spotlight + poster)

- Deploying AI in Healthcare (DAIH) Workshop, COLM 2026, San Francisco, October 2026. Non-archival.

When the Language Changes, Is It Still Safe? Multilingual Evaluation of Clinical Voice Agents | *DAIH*

Workshop, COLM 2026 (poster)

- Deploying AI in Healthcare (DAIH) Workshop, COLM 2026, San Francisco, October 2026. Non-archival.

EDUCATION

Imperial College London

2023 – 2024

MSc in Applied Machine Learning

London, United Kingdom

- Grade: Distinction.
- Relevant coursework: Machine Learning, Deep Learning (CNN, RNN, RL, autoencoders, VAE, GAN), Advanced Deep Learning Systems, Topics in Large Dimensional Data Processing, Computer Vision and Pattern Recognition, Optimization.
- Dissertation (Advisor: Dr. Sonali Parbhoo): a max-margin inverse RL method to recover the reward models of RLHF-tuned LLMs, reaching ~80% accuracy at recovering the ground-truth objective; models fine-tuned on the recovered reward showed lower toxicity than standard RLHF. Published at COLM 2025.

University of the Witwatersrand

2018 – 2021

BSc in Electrical (Information) Engineering

Johannesburg, South Africa

- Grade: Distinction; ranked 1st in the final year.
- Final-year thesis (Advisor: Dr. Vered Aharonson): Quantifying the Effectiveness of Scientific Speeches on Climate Change, marked 97%. Led to four peer-reviewed publications.

TECHNICAL SKILLS

Languages: Python, JavaScript, C++

Technologies: PyTorch, Hugging Face, LangChain, TensorFlow, Scikit-learn, Pandas, NumPy, React.js, Django

Developer Tools: Git, Docker, SQL (PostgreSQL, MySQL, NoSQL), High Performance Computing

Cloud: AWS (Lambda, S3, SageMaker), Azure (OpenAI Service, AI Studio, Cognitive Services)

ACHIEVEMENTS

- Entelect Prize, ranked 1st in final-year cohort (2021); Isazi Consulting Prize, top 5 (2020)