

Jared Joselowitz

📍 London, United Kingdom

✉ jaredjoss123@gmail.com

☎ +44 7440 408647

🌐 jaredjoselowitz

🔗 jaredjoss

SUMMARY

Lead AI Research Engineer owning the clinical safety-evaluation programme that gates every production release of LLM voice agents running across 200,000+ live patient calls. Leads a 4-person science team across safety-focused evaluation, multi-agent simulation, LLM-as-judge methods, and ASR robustness. MSc in Applied Machine Learning (Distinction), Imperial College London; first author at COLM 2025 on LLM interpretability through inverse reinforcement learning, co-author at ACL 2025 and IWSDS 2026.

EXPERIENCE

Ufonia Feb 2025 – present
Lead AI Research Engineer (promoted from Senior, Apr 2026) London, United Kingdom

- Own the clinical safety-evaluation programme that gates every production release of the LLM voice agents; set research direction for the science team and work directly with clinicians to build the validation evidence base behind release and regulatory decisions.
- Co-developed MATRIX, a multi-agent simulation framework for safety evaluation of clinical dialogue agents; its behaviour judge matched expert-clinician hazard detection (F1 0.96) in a blinded 240-dialogue assessment, across 2,100 simulated dialogues, 14 hazard scenarios, and 10 clinical domains.
- Built the ASR-safety evaluation behind WER is Unaware, a clinician-annotated benchmark and a GEPA-optimised LLM-as-judge reaching 90% accuracy and Cohen's kappa 0.816 against clinician safety assessments.
- Rebuilt prompt development around DSPy, replacing manual prompt engineering with automatic optimisation (GEPA optimiser) across both the clinical voice agents and the LLM judges to make prompts reproducible and auditable.
- Built the safety-evaluation pipelines behind these systems (reference-based, simulated-patient, and LLM-as-judge methods), spanning hazards from ASR errors to patient interruptions and multilingual interactions.

The Awareness Company Sept 2022 – Sept 2023
Data Scientist Johannesburg, South Africa

- Worked in a cross-functional team on an AI-powered data-storytelling platform, investigating and improving the robustness of production ML models.
- Designed, trained, and deployed custom ML and computer-vision models (Azure Cognitive Services): a soil-moisture predictor, a poaching-risk model for Kafue National Park (Zambia), and an energy-demand forecaster.
- Used Azure OpenAI Service to fine-tune an LLM chatbot on company data, deployed to end users via an API.

Tesseract Feb 2022 – Aug 2022
Junior Full-Stack Software Engineer Johannesburg, South Africa

- Worked in a five-person team on a full-stack AI annotation platform for business intelligence in the FinTech industry, across frontend (React) and backend (Django).
- Built end-to-end testing with Cypress and a WhatsApp annotation bot deployed on AWS Lambda.

PUBLICATIONS

Insights from the Inverse: Reconstructing LLM Training Goals Through Inverse Reinforcement Learning [arxiv.org](https://arxiv.org/abs/2505.12345)

Jared Joselowitz, Ritam Majumdar, Arjun Jagota, Matthieu Bou, Nyal Patel, Satyapriya Krishna, Sonali Parbhoo
Conference on Language Modeling (COLM 2025). [Code](#)

WER is Unaware: Assessing How ASR Errors Distort Clinical Understanding in Patient Facing Dialogue [aclanthology.org](https://aclanthology.org/2025.wer)

Zachary Ellis, Jared Joselowitz, Yash Deo, Yajie He, Anna Kalygina, Aisling Higham, Mana Rahimzadeh, Yan Jia, Ibrahim Habli, Ernest Lim
International Workshop on Spoken Dialogue Systems (IWSDS 2026), Oral Presentation. [Code](#)

ASTRID: An Automated and Scalable TRIaD for the Evaluation of RAG-based Clinical Question Answering Systems [aclanthology.org](https://aclanthology.org/2025.astrid)

Yajie Vera He, Mohita Chowdhury, Jared Joselowitz, Aisling Higham, Ernest Lim
Findings of the Association for Computational Linguistics (ACL 2025)

MATRIX: Multi-Agent simulaTion fRamework for safe Interactions and conteXtual clinical conversational evaluation [arxiv.org](https://arxiv.org/abs/2505.12345)

Ernest Lim, Yajie Vera He, Jared Joselowitz, Kate Preston, Mohita Chowdhury, Louis Williams, Aisling Higham, Katrina Mason, Mariane Melo, Tom Lawton, Yan Jia, Ibrahim Habli
arXiv preprint, under review (2025)

TALKS & PRESENTATIONS

Shipping AI to a Million Patients Without an A/B Test

June 2026

Stage talk · AI Engineer World's Fair 2026 · San Francisco

How Ufonia ships and keeps improving a regulated clinical voice agent without A/B testing on patients: simulate the highest-risk scenarios at scale, then improve it with reproducible, auditable prompt optimisation, using simulation as the release gate.

When Patients Cut In: Extending Clinical Conversational AI Safety to Interruptions

October 2026

Oral spotlight + poster · DAIH Workshop at COLM 2026 · San Francisco

Shows that when a patient interrupts a clinical voice agent mid-question, it can silently drop clinically required content that cooperative transcripts hide, and introduces a way to measure how well agents recover across different types of interruption.

When the Language Changes, Is It Still Safe? Multilingual Evaluation of Clinical Voice Agents

October 2026

Poster · DAIH Workshop at COLM 2026 · San Francisco

Shows that an agent proven safe in English can quietly degrade in other languages: simulating consultations in English, French, Polish and Urdu, emergencies are still caught but symptom-gathering and follow-up slip unevenly.

EDUCATION

Imperial College London

2023 – 2024

MSc in Applied Machine Learning

London, United Kingdom

- Grade: Distinction.
- Relevant coursework: Machine Learning, Deep Learning (CNN, RNN, RL, autoencoders, VAE, GAN), Advanced Deep Learning Systems, Topics in Large Dimensional Data Processing, Computer Vision and Pattern Recognition, Optimization.
- Dissertation (Advisor: Dr. Sonali Parbhoo): a max-margin inverse RL method to recover the reward models of RLHF-tuned LLMs, reaching ~80% accuracy at recovering the ground-truth objective; models fine-tuned on the recovered reward showed lower toxicity than standard RLHF. Published at COLM 2025.

University of the Witwatersrand

2018 – 2021

BSc in Electrical (Information) Engineering

Johannesburg, South Africa

- Grade: Distinction; ranked 1st in the final year.
- Final-year thesis (Advisor: Dr. Vered Aharonson): Quantifying the Effectiveness of Scientific Speeches on Climate Change, marked 97%. Led to four peer-reviewed publications.

TECHNICAL SKILLS

Languages: Python, JavaScript, C++

ML & LLMs: DSPy, PyTorch, Hugging Face (Transformers, Datasets, TRL, PEFT), LangChain, scikit-learn, Pandas, NumPy

LLM providers & inference: OpenAI, Anthropic, Google, Llama, OpenRouter, Ollama

Developer Tools: Git, Docker, Weights & Biases, MLflow, SQL (PostgreSQL, MySQL, NoSQL), High Performance Computing

Cloud: Google Cloud (Cloud Run, Vertex AI, BigQuery, Pub/Sub), AWS (Lambda, S3, SageMaker), Azure (AI Studio, Cognitive Services)

ACHIEVEMENTS

- Entelect Prize, ranked 1st in final-year cohort (2021); Isazi Consulting Prize, top 5 (2020)