Jared Joselowitz

London, United Kingdom | jaredjoss123@gmail.com | +44 7440 408647 | jaredjoselowitz | jaredjoss

SUMMARY

Lead AI Research Engineer owning the clinical safety-evaluation programme that gates every production release of LLM voice agents running across 200,000+ live patient calls. Leads a 4-person science team across safety-focused evaluation, multi-agent simulation, LLM-as-judge methods, and ASR robustness. MSc in Applied Machine Learning (Distinction), Imperial College London; first author at COLM 2025 on LLM interpretability through inverse reinforcement learning, co-author at ACL 2025 and IWSDS 2026.

EXPERIENCE

Ufonia — Feb 2025 – present
Lead AI Research Engineer (promoted from Senior, Apr 2026) — *London, United Kingdom*

- Own the clinical safety-evaluation programme that gates every production release of the LLM voice agents; set research direction for the science team and work directly with clinicians to build the validation evidence base behind release and regulatory decisions.
- Co-developed MATRIX, a multi-agent simulation framework for safety evaluation of clinical dialogue agents; its behaviour judge matched expert-clinician hazard detection (F1 0.96) in a blinded 240-dialogue assessment, across 2,100 simulated dialogues, 14 hazard scenarios, and 10 clinical domains.
- Built the ASR-safety evaluation behind WER is Unaware, a clinician-annotated benchmark and a GEPA-optimised LLM-as-judge reaching 90% accuracy and Cohen's kappa 0.816 against clinician safety assessments.
- Rebuilt prompt development around DSPy, replacing manual prompt engineering with automatic optimisation (GEPA optimiser) across both the clinical voice agents and the LLM judges to make prompts reproducible and auditable.
- Built the safety-evaluation pipelines behind these systems (reference-based, simulated-patient, and LLM-as-judge methods), spanning hazards from ASR errors to patient interruptions and multilingual interactions.

The Awareness Company — Sept 2022 – Sept 2023
Data Scientist — *Johannesburg, South Africa*

- Worked in a cross-functional team on an AI-powered data-storytelling platform, investigating and improving the robustness of production ML models.
- Designed, trained, and deployed custom ML and computer-vision models (Azure Cognitive Services): a soil-moisture predictor, a poaching-risk model for Kafue National Park (Zambia), and an energy-demand forecaster.
- Used Azure OpenAI Service to fine-tune an LLM chatbot on company data, deployed to end users via an API.

Tesserae — Feb 2022 – Aug 2022
Junior Full-Stack Software Engineer — *Johannesburg, South Africa*

- Worked in a five-person team on a full-stack AI annotation platform for business intelligence in the FinTech industry, across frontend (React) and backend (Django).
- Built end-to-end testing with Cypress and a WhatsApp annotation bot deployed on AWS Lambda.

PUBLICATIONS

Insights from the Inverse: Reconstructing LLM Training Goals Through Inverse Reinforcement Learning — arxiv.org
Jared Joselowitz, Ritam Majumdar, Arjun Jagota, Matthieu Bou, Nyal Patel, Satyapriya Krishna, Sonali Parbhoo
Conference on Language Modeling (COLM 2025). Code

WER is Unaware: Assessing How ASR Errors Distort Clinical Understanding in Patient Facing Dialogue — aclanthology.org
Zachary Ellis, **Jared Joselowitz**, Yash Deo, Yajie He, Anna Kalygina, Aisling Higham, Mana Rahimzadeh, Yan Jia, Ibrahim Habli, Ernest Lim
International Workshop on Spoken Dialogue Systems (IWSDS 2026), Oral Presentation. Code

ASTRID: An Automated and Scalable TRIaD for the Evaluation of RAG-based Clinical Question Answering Systems — aclanthology.org
Yajie Vera He, Mohita Chowdhury, **Jared Joselowitz**, Aisling Higham, Ernest Lim
Findings of the Association for Computational Linguistics (ACL 2025)

MATRIX: Multi-Agent simulaTion fRamework for safe Interactions and conteXtual clinical conversational evaluation — arxiv.org
Ernest Lim, Yajie Vera He, **Jared Joselowitz**, Kate Preston, Mohita Chowdhury, Louis Williams, Aisling Higham, Katrina Mason, Mariane Melo, Tom Lawton, Yan Jia, Ibrahim Habli
arXiv preprint, under review (2025)

TALKS & PRESENTATIONS

Shipping AI to a Million Patients Without an A/B Test June 2026

Stage talk · AI Engineer World's Fair 2026 · San Francisco. YouTube

How Ufonia ships and keeps improving a regulated clinical voice agent without A/B testing on patients: simulate the highest-risk scenarios at scale, then improve it with reproducible, auditable prompt optimisation, using simulation as the release gate.

When Patients Cut In: Extending Clinical Conversational AI Safety to Interruptions October 2026

Oral spotlight + poster · DAIH Workshop at COLM 2026 · San Francisco

Shows that when a patient interrupts a clinical voice agent mid-question, it can silently drop clinically required content that cooperative transcripts hide, and introduces a way to measure how well agents recover across different types of interruption.

When the Language Changes, Is It Still Safe? Multilingual Evaluation of Clinical Voice Agents October 2026

Poster · DAIH Workshop at COLM 2026 · San Francisco

Shows that an agent proven safe in English can quietly degrade in other languages: simulating consultations in English, French, Polish and Urdu, emergencies are still caught but symptom-gathering and follow-up slip unevenly.

EDUCATION

Imperial College London 2023 – 2024

MSc in Applied Machine Learning *London, United Kingdom*

- Grade: Distinction.
- Relevant coursework: Machine Learning, Deep Learning (CNN, RNN, RL, autoencoders, VAE, GAN), Advanced Deep Learning Systems, Topics in Large Dimensional Data Processing, Computer Vision and Pattern Recognition, Optimization.
- Dissertation (Advisor: Dr. Sonali Parbhoo): a max-margin inverse RL method to recover the reward models of RLHF-tuned LLMs, reaching ~80% accuracy at recovering the ground-truth objective; models fine-tuned on the recovered reward showed lower toxicity than standard RLHF. Published at COLM 2025.

University of the Witwatersrand 2018 – 2021

BSc in Electrical (Information) Engineering *Johannesburg, South Africa*

- Grade: Distinction; ranked 1st in the final year.
- Final-year thesis (Advisor: Dr. Vered Aharonson): Quantifying the Effectiveness of Scientific Speeches on Climate Change, marked 97%. Led to four peer-reviewed publications.

TECHNICAL SKILLS

Languages: Python, JavaScript, C++

ML & LLMs: DSPy, PyTorch, Hugging Face (Transformers, Datasets, TRL, PEFT), LangChain, scikit-learn, Pandas, NumPy

LLM providers & inference: OpenAI, Anthropic, Google, Llama, OpenRouter, Ollama

Developer Tools: Git, Docker, Weights & Biases, MLflow, SQL (PostgreSQL, MySQL, NoSQL), High Performance Computing

Cloud: Google Cloud (Cloud Run, Vertex AI, BigQuery, Pub/Sub), AWS (Lambda, S3, SageMaker), Azure (AI Studio, Cognitive Services)

ACHIEVEMENTS

- Entelect Prize, ranked 1st in final-year cohort (2021); Isazi Consulting Prize, top 5 (2020)